

Concurrent data structures using RCU

Rolfe Bozier

19-Oct-2014

Agenda

- Wait-free data structures
- How does the RCU mechanism work?
- A hash table with wait-free readers
- Conclusion
- References



This work is licensed under a Creative Commons Attribution 4.0 International License.

RCU – Read-Copy-Update

- What is RCU?
 - A synchronisation method that can often be used in place of a traditional reader-write lock
 - Used to protect data structures that are written and read at the same time
 - Extremely low overhead for readers – wait-free
 - Readers cannot block writers
 - Multiple writes are serialised
- What are the drawbacks then?
 - Multiple versions of the data structure can be present at any one time
 - Concurrent readers can see different [valid] versions
 - Some memory overhead in keeping older versions of the structure around (for a period)
- RCU has been used in the Linux kernel for over 10 years
- RCU requires 2 pieces of functionality:
 - Publish-subscribe mechanism
 - Grace periods




Publish-subscribe

- Consider the following stages in adding an element to a shared linked list
 - There are two readers, currently at the first node




Publish-subscribe

- We create a new node and set its next pointer to the correct destination
 - No reader can see the new node yet
 - Meanwhile, the second reader has moved to the last node



Publish-subscribe

- Finally, we *atomically* update the previous nodes next pointer to include the new node
 - The first reader follows the pointer to the new node (and then finally to the last node)



Publish-subscribe

- OK, this seems pretty simple, but...
 - Don't forget about compiler and CPU re-ordering of instructions!

Writer

```
prev = ...  
p = malloc(sizeof(node));  
p->value = 6;  
p->next = prev->next;  
rcu_assign_pointer(prev, p);
```

Add a write barrier
before setting
prev->next

Reader

```
list_for_each_entry_rcu(p, head, next)  
{  
    do_something_with(p->value);  
}
```

Add a read barrier
for each step
through the list



Publish-subscribe

- Different readers can see different views of the list at the “same time”
 - Reader 1 saw: 3, 6, 8
 - Reader 2 saw: 3, 8
- Publish semantics:
 - Writer updates the data structure atomically
 - Writer uses memory barriers to prevent store re-ordering
- Subscribe semantics
 - Readers are prepared for things to change on the fly
 - Readers use memory barriers to prevent load re-ordering




Grace periods

- Deleting an item from an RCU list is not quite so straightforward...
 - Remember that readers can be looking at any node in the list while we are deleting
 - ... including the node we want to remove!
- To allow for this, we preserve the existing data for a grace period, to allow readers to complete
 - After the grace period has expired, we can delete the old node
 - Obviously we need something a bit more rigorous than a “timeout”




Grace periods

- Consider the following stages in deleting an element in a shared linked list




Grace periods

- The next pointer in the previous node is updated to bypass the target node
 - But before this, reader 2 has traversed to the target node




Grace periods

- The writer waits for a “grace period”
- After the grace period has expired, there can be no readers still referencing the target node
 - Reader 2 has moved to the final node



Grace periods

- The target node is freed
 - Meanwhile the first reader traverses the list to the final node



Grace periods

- Different readers can see different views of the list at the “same time”
 - Reader 1 saw: 3, 8
 - Reader 2 saw: 3, 6, 8
- The grace period allows the writer to safely perform garbage collection of the deleted node
 - This needs to be long enough for all readers to no longer reference the deleted data
 - It should not be so long as to hold up the writer or cause excess memory usage
- If we want to *update* a node, we generally can’t do this atomically (unless the data is atomic)
 - **Read** the node to be replaced
 - **Copy** it to a new node with the modified values
 - **Update** the list atomically to use the replacement node



Grace periods

- In the Linux kernel, grace periods are managed as follows:
 - Readers enclose their operations in `rcu_read_lock()` and `rcu_read_unlock()` “functions” - in the simple case they do nothing
 - Code must *not* block between these two functions
 - The write thread schedules itself to run on each CPU in turn
 - At the end, it is guaranteed that no reader is still holding on to a deleted node
 - The node can be freed
- Obviously something different needs to be done for userspace implementations of RCU
- There are now a number of userspace libraries that can manage RCU operations in various ways, e.g.
 - A global 64-bit grace period counter that can be atomically incremented, plus
 - Per-thread grace period counter references to indicate where the reader is
- See the references for more information




An RCU Hash Table

- Consider the example of an RCU Hash Table
 - In open hashing, collisions are resolved by adding a linked list to the bucket
 - We can see that RCU lets us add, modify and delete entries to a hash table, including the bucket lists
- But what about resizing?
 - We can resize a hash table without blocking readers with the following conditions:
 - The table grows or shrinks by a factor of two
 - Bucket lists can be *imprecise* – they can contain extra entries (from different hash values)



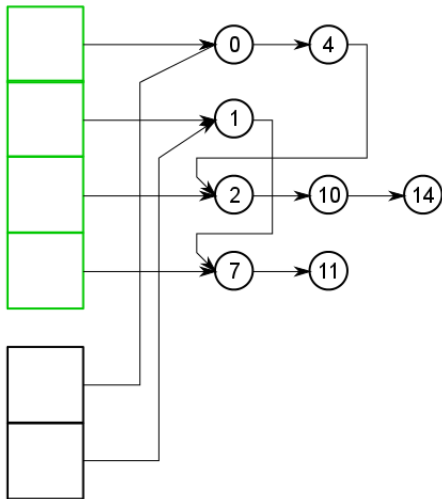
An RCU Hash Table

- Consider a hash table with 4 buckets
- We want to shrink it to size 2



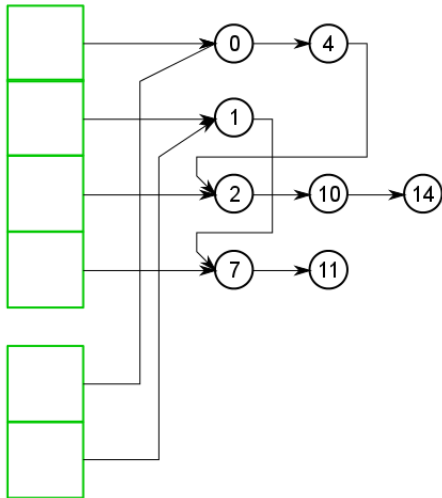
An RCU Hash Table

- Create the new bucket array
- Chain the bucket lists according to their new hash value



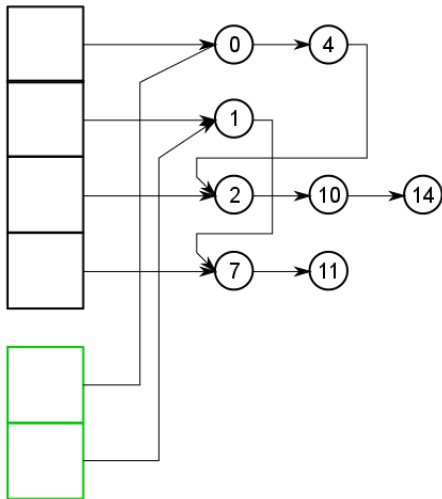
An RCU Hash Table

- Publish the new table and wait for readers



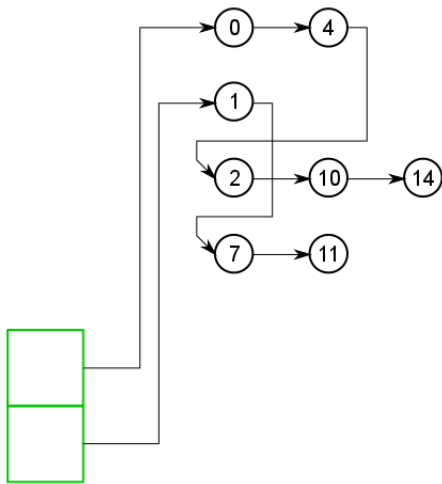
An RCU Hash Table

- The old bucket array can now be freed



An RCU Hash Table

- The final version



- Growing is essentially the reverse process
 - It is a bit more involved due to the need to unzip the bucket lists and wait for readers for each step

Conclusion

- RCU is a synchronisation mechanism for managing data structures with wait-free readers
 - Most suitable where reads \gg writes
 - The data structure semantics may need some adjustment (e.g. imprecise bucket lists)
- It has been used in Linux kernels for over 10 years
- To implement, you need a few things
 - Atomic operations
 - Memory barriers
 - A suitable grace period mechanism



References

- What is RCU, Fundamentally? [3 parts]
 - <http://lwn.net/Articles/262464/>
 - <http://lwn.net/Articles/263130/>
 - <http://lwn.net/Articles/264090/>
- “User-Level Implementations of Read-Copy Update”
 - <https://www.efficios.com/pub/rcu/urcu-main.pdf>
 - <https://www.efficios.com/pub/rcu/urcu-suppl.pdf>
- Resizable, Scalable, Concurrent Hash Tables via Relativistic Programming
 - http://www.usenix.org/events/atc11/tech/final_files/Triplett.pdf
 - http://works.bepress.com/jonathan_walpole/12/

